



# Database Management Systems


An introduction to DMS

Duration: 30 minutes

Q&A: 5 minutes by the end of the lecture

## In our previous app:

We used the Node.js **fs** module to store data in files.



## Our persistence solution should

- Only rewrite portions of state on change
- Be capable of storing more than can be held in memory
- Efficiently access and address the stored information in multiple, flexible, useful ways
- Allow more than one process simultaneous access
- Avoid inconsistency of state due to changes
- Retain information through a power failure

## Our persistence solution should


- Only rewrite portions of state on change
- Be capable of storing more than can be held in memory
- ✗ Efficiently access and address the stored information in multiple, flexible, useful ways
- ✗ Allow more than one process simultaneous access
- ✗ Avoid inconsistency of state due to changes
- Retain information through a power failure

## Our persistence solution should

- ❌ Efficiently access and address the stored information in multiple, flexible, useful ways
- ❌ Allow more than one process simultaneous access
- ❌ Avoid inconsistency of state due to changes

*What is holding our solution back?*

**The data is disorganized!**



## Looking forward...

Can we find an alternative non-volatile storage mechanism?

- Better organizes the data
- Increases the efficiency of operations on the data

# Databases



# Database

- An **organized** collection of related data

What are some good ways to **organize** data ?

Data Structures!

# Databases

Just like data structures, databases use similar primitives:

- Key-Value pairs
- Tables
- Plain text documents
- Graphs

# Databases

Though different, all databases organize data such that it is predictable and highly functional

**Data Table**

id	user	text	room
412	fred	aw!	lobby
413	sunny	snap	park
414	allen	omg	park
415	pamela	yo!	lobby
416	fred	dude	park
417	pamela	huh?	park
418	sunny	umm	lobby
419	sunny	yo!	roof
420	fred	duh!	park
421	allen	snap	lobby

**Text Document**

```
{'id':412,'user':'fred','text':'aw!','room':'lobby'}|{'id':413,'user':'sunny','text':'snap','room':'park'}|{'id':414,'user':'allen','text':'omg','room':'park'}|{'id':415,'user':'pamela','text':'yo!','room':'lobby'}|{'id':416,'user':'fred','text':'dude','room':'park'}|{'id':417,'user':'pamela','text':'huh?','room':'park'}|{'id':418,'user':'sunny','text':'umm','room':'lobby'}|{'id':419,'user':'sunny','text':'yo!','room':'roof'}|{'id':420,'user':'fred','text':'duh!','room':'park'}|{'id':421,'user':'allen','text':'snap','room':'lobby'}
```

# Databases

Though different, all databases organize data such that it is predictable and highly functional

**Data Table**

id	user	text	room
412	fred	aw!	lobby
413	sunny	snap	park
414	allen	omg	park
415	pamela	yo!	lobby
416	fred	dude	park
417	pamela	huh?	park
418	sunny	umm	lobby
419	sunny	yo!	roof
420	fred	duh!	park
421	allen	snap	lobby

**Text Document**

```
{'id':412,'user':'fred','text':'aw!','room':'lobby'}|{'id':413,'user':'sunny','text':'snap','room':'park'}|{'id':414,'user':'allen','text':'omg','room':'park'}|{'id':415,'user':'pamela','text':'yo!','room':'lobby'}|{'id':416,'user':'fred','text':'dude','room':'park'}|{'id':417,'user':'pamela','text':'huh?','room':'park'}|{'id':418,'user':'sunny','text':'umm','room':'lobby'}|{'id':419,'user':'sunny','text':'yo!','room':'roof'}|{'id':420,'user':'fred','text':'duh!','room':'park'}|{'id':421,'user':'allen','text':'snap','room':'lobby'}
```

Note that data is usually stored in different and more efficient encodings like BSON

# Databases

For now, let's focus on **Relational Databases**.

# Databases

In Relational Databases, data is stored in tables

- **Each table represents a data entity**
- The table columns represent attributes of the entity
- Each row in the table represents a single record or instance

Messages			
ID	User	Text	Room
412	fred	aw!	lobby
413	sunny	snap	park
414	allen	omg	park
415	pamela	yo!	lobby
416	fred	dude	park
417	pamela	huh	park
418	sunny	umm	lobby

# Databases

In Relational Databases, data is stored in tables

- Each table represents a data entity
- **The table columns represent attributes of the entity**
- Each row in the table represents a single record or instance

Messages			
ID	User	Text	Room
412	fred	aw!	lobby
413	sunny	snap	park
414	allen	omg	park
415	pamela	yo!	lobby
416	fred	dude	park
417	pamela	huh	park
418	sunny	umm	lobby

# Databases

In Relational Databases, data is stored in tables

- Each table represents a data entity
- The table columns represent attributes of the entity
- **Each row in the table represents a single record or instance**

Messages			
ID	User	Text	Room
412	fred	aw!	lobby
<b>413</b>	<b>sunny</b>	<b>snap</b>	<b>park</b>
414	allen	omg	park
415	pamela	yo!	lobby
416	fred	dude	park
417	pamela	huh	park
418	sunny	umm	lobby



# Relational Databases

Relational Database tables have hard constraints

- The size and type of each column is fixed
- Data that does not fit these constraints will be truncated/ rejected

```
| id | user      | text      | room
|
| ---| -|-----| -|-----| -|-----
|
| 412| fred---| aw!----| lobby-
|
| 413| sunny--| snap---| park--
|
| 414| allen--| omg----| park--
|
```

# Databases

Relational Database tables have hard constraints

- The size and type of each column is fixed
- Data that does not fit these constraints will be truncated/ rejected

id	user	text	room
---	-----	-----	-----
412	fred---	aw!----	lobby-
413	sunny--	snap---	park--
414	allen--	omg----	park--

Could the message 'Greetings' from user 'Samantha' in room 'Kitchen' be added?

# Relational Databases

Tabular constraints make the data easier to work with.

- The size of each record is predefined and predictable
- This predictability improves the performance accessing information.

```
|id |user  |text  |room
|---|-----|-----|-----
|412|fred  |Fixed width |lobby|
|413|sunny--|snap---|park--|
|414|allen--|omg----|park--|
```

# Relational Databases

Tabular constraints make the data easier to work with.


- The size of each record is predefined and predictable
- This predictability improves the performance of accessing information.
- SQL DBs uses array-like techniques to achieve  **$O(1)$**

```
...|412|fred|Fixed width|---|lobby|413|sun|Fixed width|---|park|414|alle
n--|omg---|park--|415|pa|Fixed width|---|lobby|416|f|Fixed width|do---|
park--|417|pamela-|huh?---|park--|418|sunny--|umm---|lobby|419|s
unny--|yo!----|roof--|420|fred---|duh!---|park--
--|lobby|422...
```

Fixed width data format  
enables array-like  
constant-time access to any  
index.

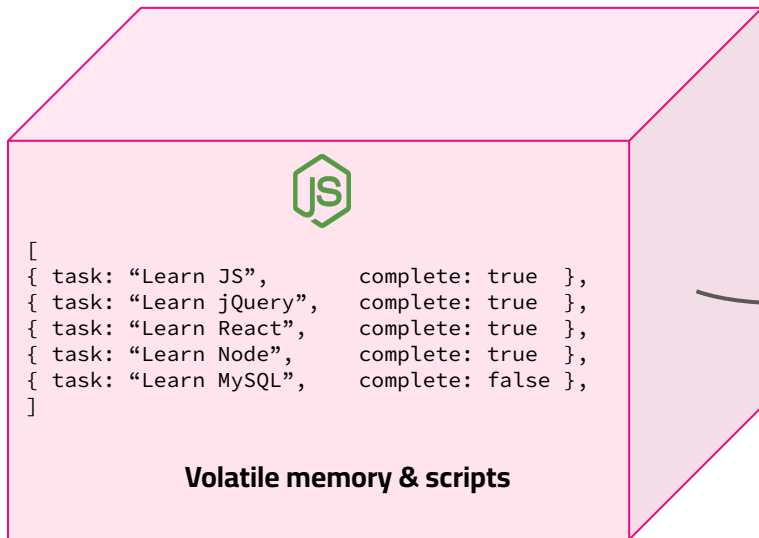
# Relational Databases

We can store reliably organized, efficient data structures.




# Relational Databases

How do we create/ manage/ interact with this data ?




?



# Relational Databases

How do we create/ manage/ interact with this data ?



# Database Management Systems



# Database Management Systems

The software that optimizes and manages the storage and retrieval of data from databases.

- A database is just a collection of data
- A database management system enables you to interact with that database

\*Often the term “ database” is used to describe both the data and the DBMS.

# Database Management Systems



# Database Management Systems

DBMS's achieve similar goals through different means

- MySQL and SQLite both maintain relational databases
  - SQLite is server-less and self-contained, creating databases which exist entirely within the project files
  - MySQL runs on an isolated server allowing for many advanced features and greater performance

# Database Management Systems

How does this fit into my application?

- A DBMS is completely independent process
- To interact with your database, you must connect and communicate with the DBMS process



# Database Management Systems

How does this fit into my application?

- A DBMS is completely independent process
- To interact with your database, you must connect and communicate with the DBMS process



The DBMS becomes the third separate layer in our stack

# Connecting to your DBMS

What is needed to communicate between two processes?

- A connection
  - We must be able to connect from some client
- A communication protocol
  - Probably not HTTP in this case
- Messages
  - In a language the recipient understands

# Connecting to your DBMS

What is needed to communicate between two processes?

- **A connection**
  - **We must be able to connect from some client**
- A communication protocol
  - Probably not HTTP in this case
- Messages
  - In a language the recipient understands

# Connecting to your DBMS

## Creating a DBMS connection

- To connect to the DBMS from our Node.js server, we can install a driver module.
- Drivers expose a JS interface and transmit commands.
- Could we connect clients other than our server?



# Connecting to your DBMS

What is needed to communicate between two processes?

- A connection
  - We must be able to connect from some client
- **A communication protocol**
  - **Probably not HTTP in this case**
- Messages
  - In a language the recipient understands

Thankfully, we don't typically worry about this since it's hidden behind a JS interface

# Connecting to your DBMS

What is needed to communicate between two processes?

- A connection
  - We must be able to connect from some client
- A communication protocol
  - Probably not HTTP in this case
- **Messages**
  - **In a language the recipient understands**

# Connecting to your DBMS

Database Management Systems don't speak JavaScript

- JavaScript is great scripting language, but it isn't great for data access and manipulation
- DBMS's will use specialized, query languages
- These languages will focus on individual statements that can be used to query the database
- Languages often resemble the shape of the data within the databases they access

# Query Languages

Database Management Systems don't speak JavaScript

- DBMS's will use specialized, query languages
- Languages often resemble the shape of the data within the databases they access
  - e.g Cypher is used for the graph database Neo4J

```
MATCH (user: {name: 'John' } )-[:FRIEND]->(follower)
```

# Query Languages

Database Management Systems don't speak JavaScript

- DBMS's will use specialized, query languages
- Languages often resemble the shape of the data within the databases they access
  - e.g Cypher is used for the graph database Neo4J
  - **SQL is used for relational databases**

```
SELECT * FROM friends f JOIN users u ON u.id = f.id;
```

# Structured Query Language

```
mysql> show tables;
```

First, let's look at all the tables in our database with **show tables**.

```
mysql> show tables;
```

```
+-----+  
| Tables_in_school |  
+-----+  
| classes           |  
| classes_students  |  
| students          |  
| teachers          |  
+-----+  
4 rows in set (0.00 sec)
```

First, let's look at all the tables in our database with **show tables**.



```
mysql> describe classes;
```

We can use the **describe** command to look at the columns for a particular table.

```
mysql> describe classes;
```

Field	Type	Null	Key	Default	Extra
id	int(11)	NO	PRI	NULL	auto_increment
name	varchar(225)	NO		NULL	
room	varchar(255)	NO		NULL	
teacher_id	int(11)	YES	MUL	NULL	

4 rows in set (0.04 sec)

We can use the **describe** command to look at the columns for a particular table.

```
mysql> select * from classes;
```

Now let's look at the specific data in our classes table.

```
mysql> select * from classes;
```

id	name	romm	techer_id
1	CS 101	114	2
2	CS 102	114	2
3	CS 151	222	3
4	CS 245	118	3
5	CS 330	220	3
6	PSY101	318	1
7	PSY201	318	1

```
7 rows in set (0.00 sec)
```

The asterisk character tells MySQL “I want to see data for all of the columns in this table”.

```
mysql> select name, room from classes;
```

What if we only wanted to see the class name and its location? We can specify columns by referring to them by name in our query.

```
mysql> select name, room from classes;
```

name	room
CS 101	114
CS 102	114
CS 151	222
CS 245	118
CS 330	220
PSY101	318
PSY201	318

```
7 rows in set (0.01 sec)
```

What if we only wanted to see the class name and its location? We can specify columns by referring to them by name in our query.

```
mysql> select name, room from classes where teacher_id=2;
```

Let's start imagining some real world scenarios. Let's say I know a particular teacher's ID number and would like to see their class schedule.

```
mysql> select name, room from classes where teacher_id=2;
```

```
+-----+-----+  
| name   | room |  
+-----+-----+  
| CS 101 | 114   |  
| CS 102 | 114   |  
+-----+-----+  
2 rows in set (0.01 sec)
```

This is easy because we specified a `teacher_id` column in our `classes` table. The `where` keyword in our query introduces a constraint: we only want to see records that satisfy a certain condition.



```
mysql> insert into classes (name,room, teacher_id)
      values ("CS101", 25, 6);
```

Query OK, 1 row affected (0.02 sec)

Here's an example of an INSERT statement which will add values to the specified fields (columns) in a new record (row)

## Data Storage Goals, revisited

- Only rewrite portions of state on change
- Be capable of storing more than can be held in memory
- Efficiently access and address the stored information in multiple, flexible, useful ways
- Allow more than one process simultaneous access
- Avoid inconsistency of state due to changes
- Retain information through a power failure

Databases offer an ideal solution for data persistence and management

For any Full Stack application, utilizing a database for data storage is almost mandatory.

A solid pink horizontal bar with rounded ends, located in the top right corner of the slide.

# That's it